

MP5 Design

Clarification: I finish the basic requirement and also Option1 and Option2.

First, I design a FIFO scheduler for kernel level threads. To be more specific, I define a double linklist structure to support the ready queue, and also add head, tail and size attributes in class Scheduler. So with this ready queue, I can just pick up the first thread and put the switch-out thread at the tail in the ready queue when do FIFO scheduling. Also, when the thread returns, finishes, I will do the corresponding cleaning work and scheduling to other threads in the ready queue. Plus, I will also discuss a problem I encountered when the ready queue is empty and how to deal with it.

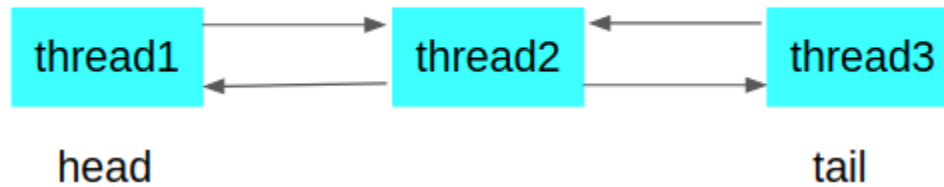
Second, for Option1, in order to support interrupt, I enable interrupts in thread_start, because the interrupt is disabled during initialization; and also enable interrupts / disable interrupts to protect the critical sections, like the add, resume, yield, terminate in Scheduler.

Third, for Option2, to implement RR Scheduling, based on the SimpleTimer, whenever 50 ms passes, I will schedule the next thread in the ready queue. Since I will perform context switching inside the SimpleTimer::handle_interrupt(), before this handle_interrupt() performs, I will send EOI message first.

Part 1 FIFO scheduler, basic requirement

(1) Define a ready queue in Scheduler, based on the double link list.

ready queue



```
57 // Double linklist structure to support the ready queue
58 typedef struct Node {
59     Thread *t;
60     Node *prev;
61     Node *next;
62 }Node;
63
64 class Scheduler {
65
66     /* The scheduler may need private members... */
67     Node *_head;
68     Node *_tail;
69     unsigned int _size;
```

(2) Main method implementation of Scheduler

Constructor:

```
Scheduler::Scheduler():_head(NULL),_tail(NULL),_size(0) {
    Console::puts("Constructed Scheduler.\n");
}
```

yield:

I use `disble_interrups` and `enable_interrups` to protect the Critical Section, notice this is a single-thread system.

If there are threads in the ready queue, pick the head and switch to it, FIFO scheduling.

Also update the ready queue structure.

```

52 void Scheduler::yield() {
53
54     if (Machine::interrupts_enabled())
55         Machine::disable_interrupts();
56
57     if (_size) {
58
59         Thread *t = _head->t;
60         Node *tmp = _head;
61         if (_size == 1) {
62             _head = NULL;
63             _tail = NULL;
64         } else {
65             _head = _head->next;
66             _head->prev = NULL;
67         }
68         --_size;
69         delete tmp;
70         Thread::dispatch_to(t);
71         //Console::puts("inside yield() after dispatch_to\n");
72     }
73
74     if (!Machine::interrupts_enabled())
75         Machine::enable_interrupts();
76
77 }

```

add:

Just add a new thread node at the tail of the ready queue.

```

83 void Scheduler::add(Thread * _thread) {
84     assert(_thread != NULL);
85
86     if (Machine::interrupts_enabled())
87         Machine::disable_interrupts();
88
89     Node *n = new Node();
90     if (!n) {
91         Console::puts("new Node() failed
92         return;
93     }
94     n->t = _thread;
95
96     if (_size) {
97         if (_size == 1) {
98             n->prev = _head;
99             n->next = NULL;
100             _head->next = n;
101             _tail = n;
102         }else {
103             n->next = NULL;
104             n->prev = _tail;
105             _tail->next = n;
106             _tail = n;
107         }
108     }else {
109         n->prev = NULL;
110         n->next = NULL;
111         _head = n;
112         _tail = n;
113     }
114
115     ++_size;
116
117     if (!Machine::interrupts_enabled())
118         Machine::enable_interrupts();
119
120 }

```

resume:

I just call add.

```
79 void Scheduler::resume(Thread * _thread) {
80     add(_thread);
81 }
```

terminate:

If the to-terminate thread is in the ready queue, delete it from the queue.

```
122 void Scheduler::terminate(Thread * _thread) {
123
124     if (Machine::interrupts_enabled())
125         Machine::disable_interrupts();
126
127
128     if (_thread && _size) {
129         Node * current = _head;
130         while (current) {
131             if (current->t == _thread) break;
132             current = current->next;
133         }
134
135         if (current) {
136
137             if (_size == 1) {
138                 delete current;
139                 _head = NULL;
140                 _tail = NULL;
141             } else if (_size == 2) {
142                 if (current == _head) {
143                     _head = _tail;
144                     _head->prev = NULL;
145                 } else { // current == _tail
146                     _tail = _head;
147                     _head->next = NULL;
148                 }
149             }
150         }
151     }
152 }
```

```

147         _head->next = NULL;
148     }
149     delete current;
150 }else {
151     if (current == _head) {
152         _head = _head->next;
153         _head->prev = NULL;
154     }else if (current == _tail) {
155         _tail = _tail->prev;
156         _tail->next = NULL;
157     }else {
158         current->prev->next = current->next;
159         current->next->prev = current->prev;
160     }
161     delete current;
162 }
163 --_size;
164 } // end if (current)
165 }
166
167 if (!Machine::interrupts_enabled())
168     Machine::enable_interrupts();
169
170 }

```

Thread::thread_shutdown

In order to deal with thread-return correctly, I call Scheduler::terminate(), delete the current_thread, and Scheduler::yield() in the Thread::thread_shutdown. Because thread_shutdown() will be called when the thread returns from the thread function, I do the corresponding cleaning work and also schedule the next thread here.

```

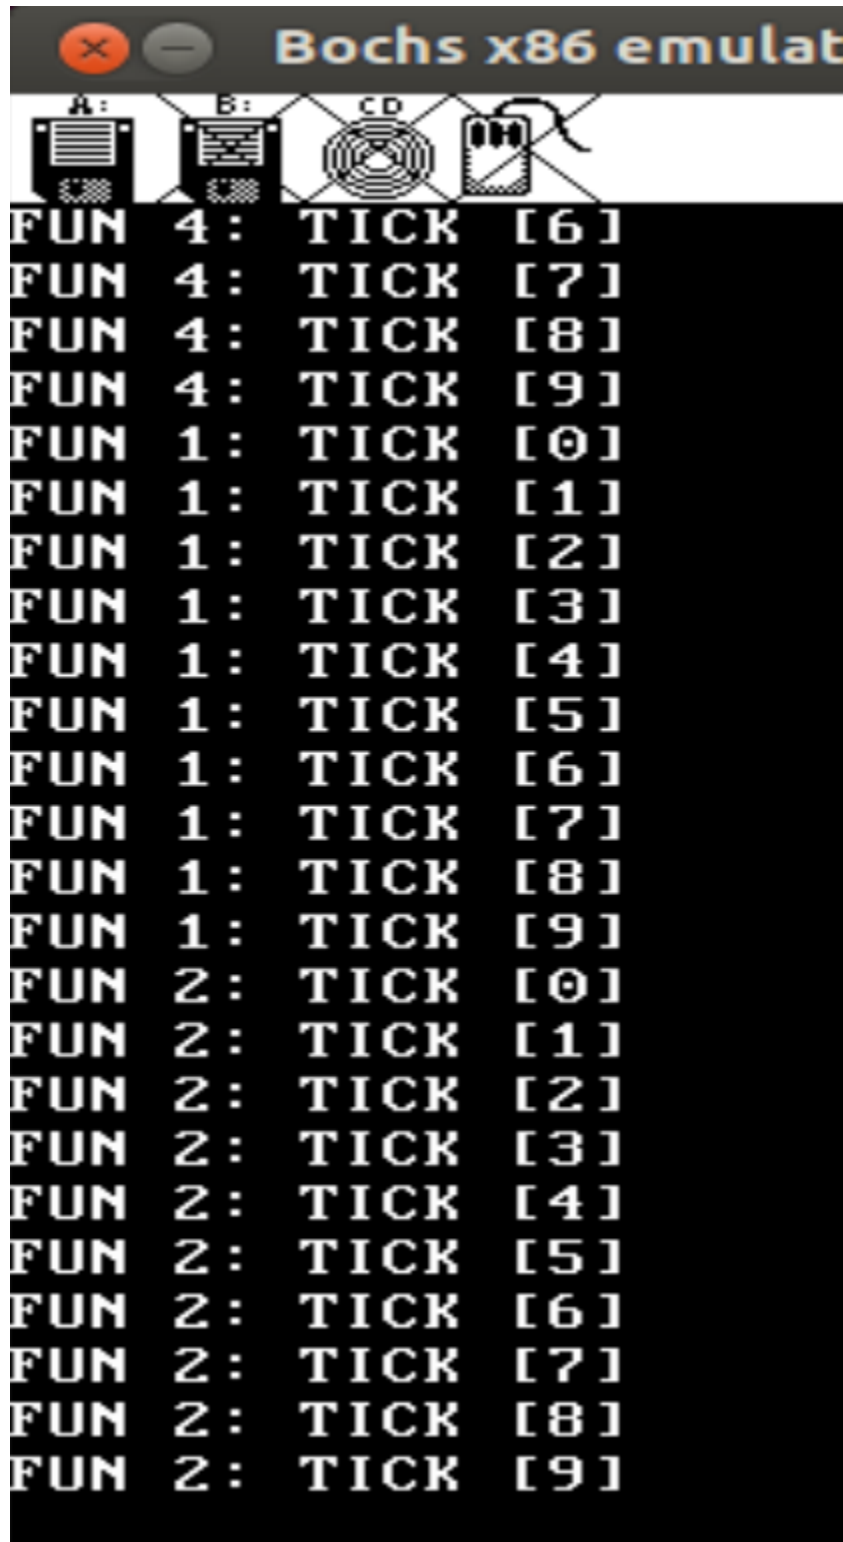
74 static void thread_shutdown() {
75     /* This function should be called when the thread
76        function.
77        It terminates the thread by releasing memory
78        held by the thread.
79        This is a bit complicated because the thread
80        has the scheduler.
81        */
82     //assert(false);
83     /* Let's not worry about it for now.
84        This means that we should have non-terminating
85        */
86     SYSTEM_SCHEDULER->terminate(current_thread);
87     delete current_thread;
88     SYSTEM_SCHEDULER->yield();
89 }
90

```

Testing for FIFO Scheduler, just the basic requirement of this MP.


(1)With #define _USES_SCHEDULER_

As shown below, Thread1 ~ Thread4 are running one by one in FIFO Scheduling. It is correct.



(2) With both `#define _USES_SCHEDULER_` and `#define _TERMINATING_FUNCTIONS_`

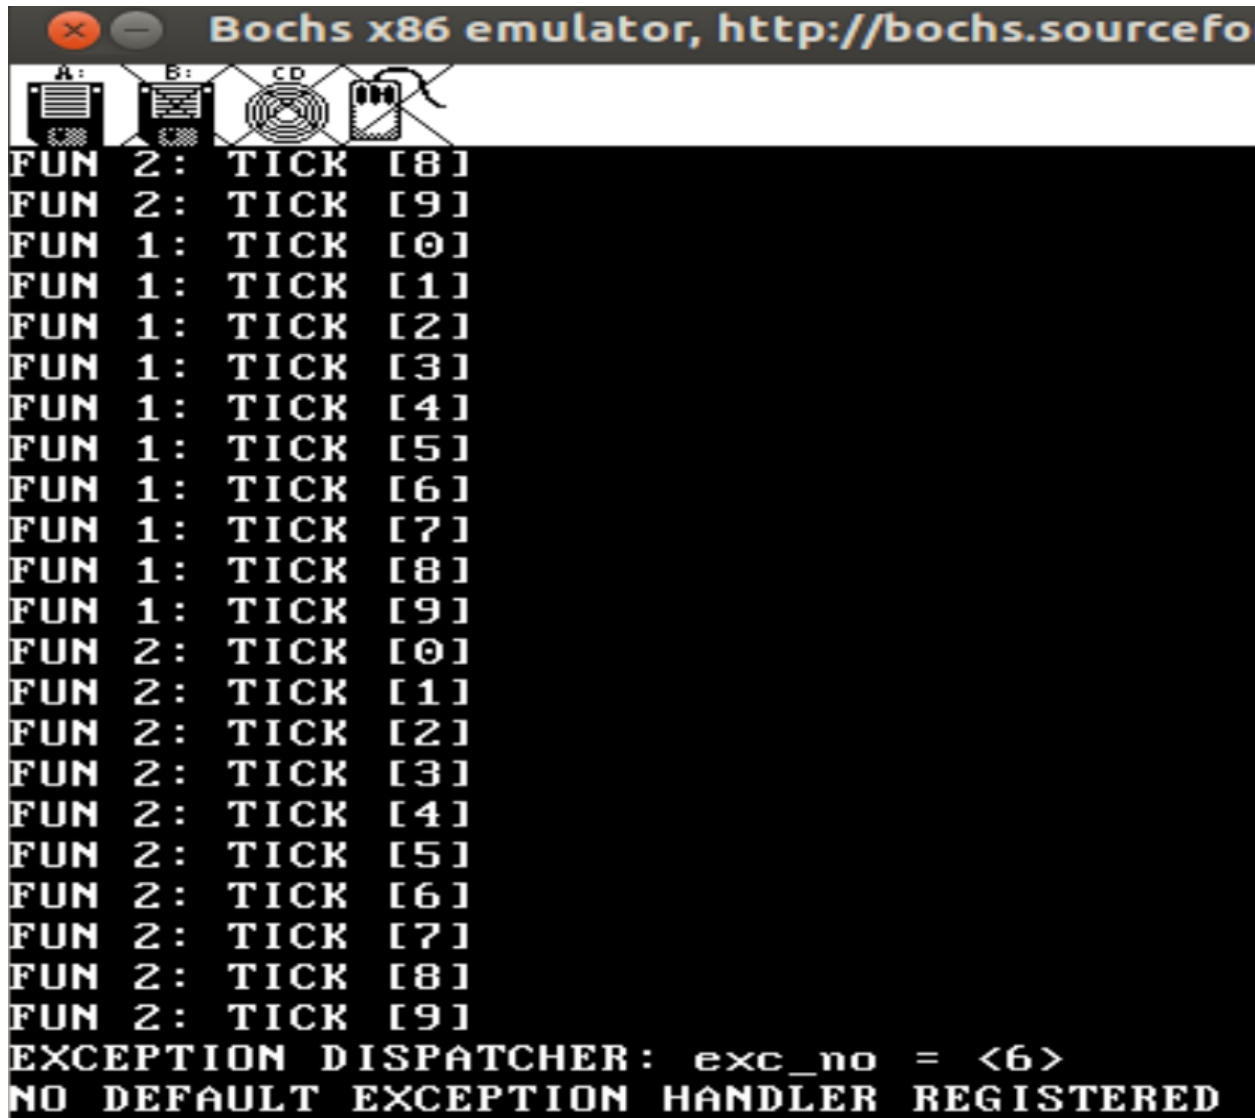
At first, Thread1 - Thread4 will run one by one just as in Testing case (1), but since Thread1 and Thread2 will terminate after 10 loops, Thread3 and Thread4 will run one by one later. This is also correct because Thread1 and Thread2 terminate successfully.



```
FUN 4: TICK [8]
FUN 4: TICK [9]
FUN 3 IN BURST[1142]
FUN 3: TICK [0]
FUN 3: TICK [1]
FUN 3: TICK [2]
FUN 3: TICK [3]
FUN 3: TICK [4]
FUN 3: TICK [5]
FUN 3: TICK [6]
FUN 3: TICK [7]
FUN 3: TICK [8]
FUN 3: TICK [8]
FUN 3: TICK [9]
FUN 4 IN BURST[1142]
FUN 4: TICK [0]
FUN 4: TICK [1]
FUN 4: TICK [2]
FUN 4: TICK [3]
FUN 4: TICK [4]
FUN 4: TICK [5]
FUN 4: TICK [6]
FUN 4: TICK [7]
FUN 4: TICK [8]
FUN 4: TICK [9]
```


A problem to discuss: What happens when the ready queue is empty?

Without infinite thread, the ready queue of the Scheduler will be empty, the system output will be like this:



```
Bochs x86 emulator, http://bochs.sourceforge...  
A: B: CD  
FUN 2: TICK [8]  
FUN 2: TICK [9]  
FUN 1: TICK [0]  
FUN 1: TICK [1]  
FUN 1: TICK [2]  
FUN 1: TICK [3]  
FUN 1: TICK [4]  
FUN 1: TICK [5]  
FUN 1: TICK [6]  
FUN 1: TICK [7]  
FUN 1: TICK [8]  
FUN 1: TICK [9]  
FUN 2: TICK [0]  
FUN 2: TICK [1]  
FUN 2: TICK [2]  
FUN 2: TICK [3]  
FUN 2: TICK [4]  
FUN 2: TICK [5]  
FUN 2: TICK [6]  
FUN 2: TICK [7]  
FUN 2: TICK [8]  
FUN 2: TICK [9]  
EXCEPTION DISPATCHER: exc_no = <6>  
NO DEFAULT EXCEPTION HANDLER REGISTERED
```

This is expected, because no thread any more in the ready queue after finishing the current thread, so it is an invalid opcode for the system. In order to prevent the OS from this exception when there is no infinite threads, we can add a dummy thread to run infinitely in the kernel.C, since there will be always at least one thread in the ready queue, avoiding the 'invalid opcode' problem.

```
42 /* UNCOMMENT THE FOLLOWING LINE IF YOU WANT A DUMMY THREAD TO RUN */  
43 #define _DUMMYTHREAD_
```


Part 2 Interrupt support, Option 1

In order to support interrupts, first I enable interrupts inside `Thread::thread_start`, since during the thread initialization, thread's interrupts will be disabled. I also disable interrupts before entering Critical Section, like adding the ready queue in `Scheduler::run`, dequeue in `Scheduler::yield()`, and enable interrupts after finishing the Critical Section.

Testing

As shown below, the timer interrupts will still occur. It is correct.



Part 3 RR Scheduling, Option 2

I support both RR Scheduling and voluntarily yielding CPU here. In order to implement 50ms quantum time of RR Scheduling, I do scheduling inside `SimpleTimer::handle_interrupt()` when 50ms passes.

```

56 void SimpleTimer::handle_interrupt(REGS *_r) {
57 /* What to do when timer interrupt occurs? In this case, we update "ticks",
58    and maybe update "seconds".
59    This must be installed as the interrupt handler for the timer in the
60    when the system gets initialized. (e.g. in "kernel.C") */
61
62    /* Increment our "ticks" count */
63    ticks++;
64
65    /* Whenever ticks 5 times, 50 ms passed */
66    if (ticks%5 == 0) {
67        Console::puts(" 50ms time quantum ends\n");
68        SYSTEM_SCHEDULER->resume(Thread::CurrentThread());
69        SYSTEM_SCHEDULER->yield();
70    }
71
72
73    /* Whenever a second is over, we update counter accordingly. */
74    if (ticks >= hz ) {
75        seconds++;
76        ticks = 0;
77        Console::puts(" One second has passed\n");
78    }
79 }

```

Since the thread will yield CPU inside the Timer interrupt handler, we need to send EOI message before performing the scheduling in InterruptHandler::dispatch_interrupt().

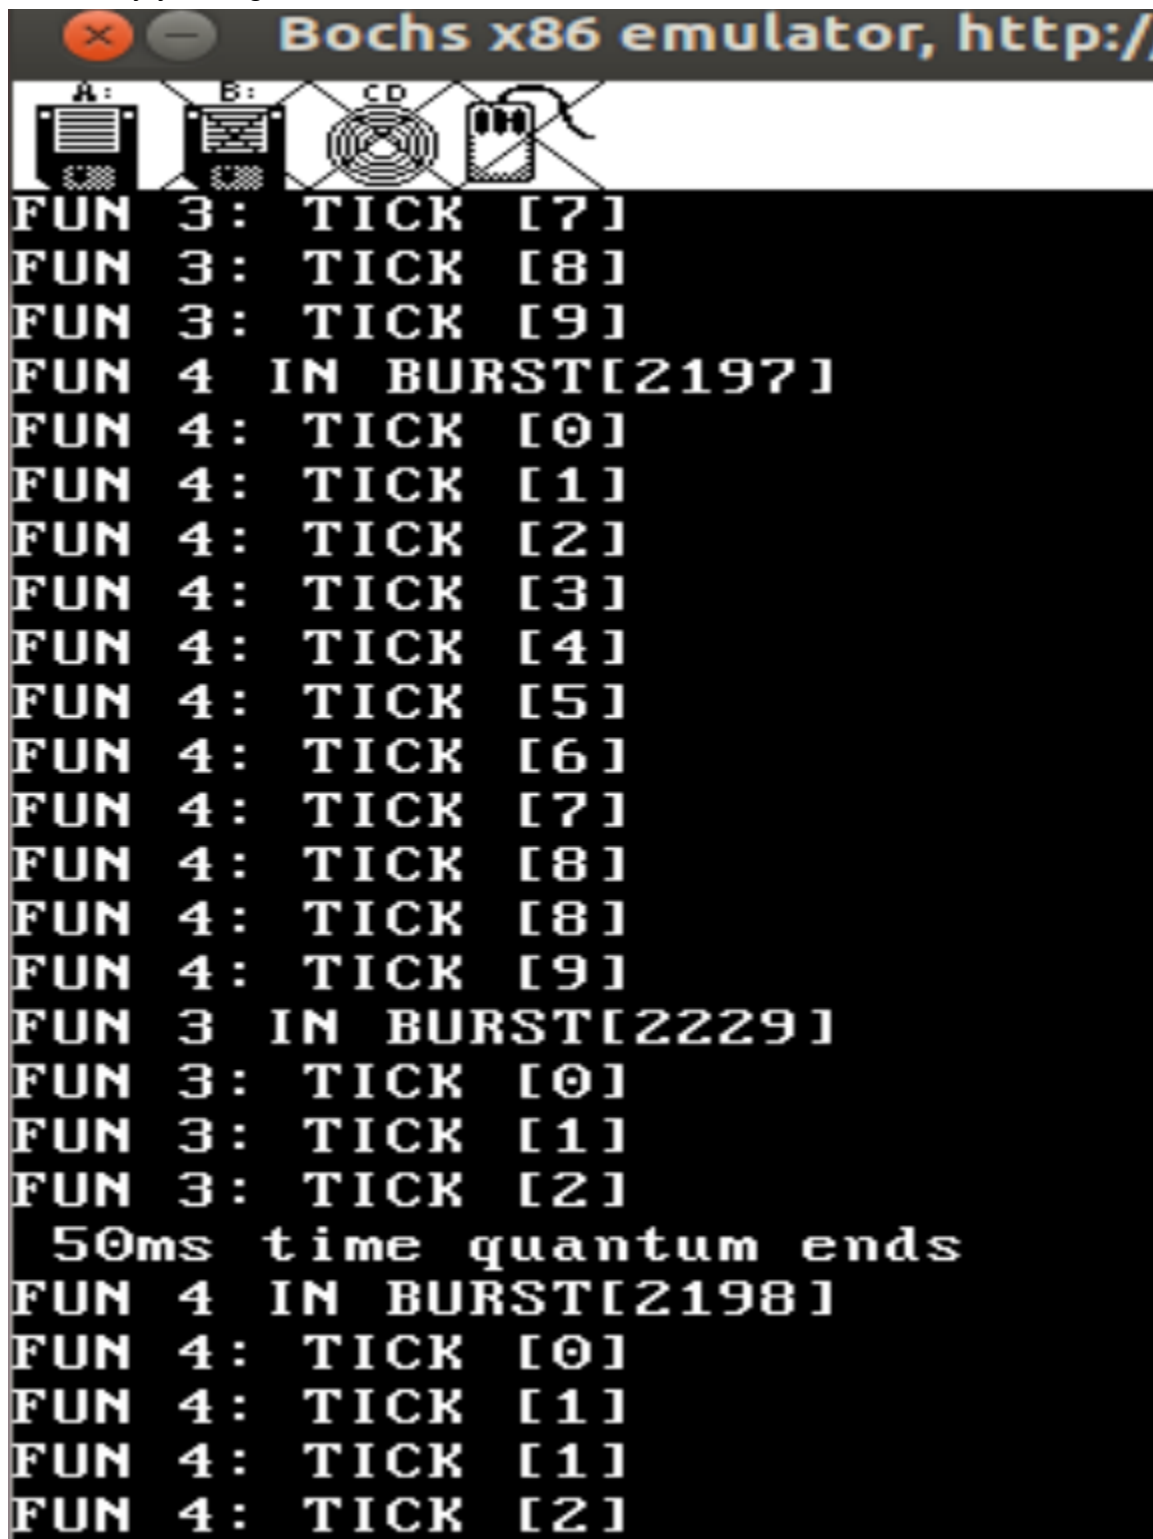
```

106 void InterruptHandler::dispatch_interrupt(REGS *_r) {
107
108    /* -- INTERRUPT NUMBER */
109    unsigned int int_no = _r->int_no - IRQ_BASE;
110
111    //Console::puts("INTERRUPT DISPATCHER: int_no = ");
112    //Console::putui(int_no);
113    //Console::puts("\n");
114
115    assert((int_no >= 0) && (int_no < IRQ_TABLE_SIZE));
116
117    /* -- HAS A HANDLER BEEN REGISTERED FOR THIS INTERRUPT NO? */
118
119    InterruptHandler * handler = handler_table[int_no];
120
121    if (!handler) {
122        /* --- NO DEFAULT HANDLER HAS BEEN REGISTERED. SIMPLY RETURN AN ERROR.
123        Console::puts("INTERRUPT NO: ");
124        Console::puti(int_no);
125        Console::puts("\n");
126        Console::puts("NO DEFAULT INTERRUPT HANDLER REGISTERED\n");
127        // abort();
128    }else {
129
130        /* This is an interrupt that was raised by the interrupt controller. We need
131        to send an end-of-interrupt (EOI) signal to the controller after the
132        interrupt has been handled. */
133
134        /* Check if the interrupt was generated by the slave interrupt controller
135        If so, send an End-of-Interrupt (EOI) message to the slave controller.
136
137        if (generated_by_slave_PIC(int_no)) {
138            Machine::outportb(0xA0, 0x20);
139        }
140
141        /* Send an EOI message to the master interrupt controller. */
142        Machine::outportb(0x20, 0x20);
143
144        /* -- HANDLE THE INTERRUPT */
145        handler->handle_interrupt(_r);
146    }
147 }

```

Testing RR Scheduling

As shown in the 3 screenshots below, my OS support RR Scheduling and also voluntarily yielding CPU.



The image shows a screenshot of the Bochs x86 emulator window. The title bar reads "Bochs x86 emulator, http://". Below the title bar is a toolbar with icons for A:, B:, CD, and a floppy disk. The main area is a black terminal window with white text. The text shows a sequence of function ticks and bursts. It starts with "FUN 3: TICK [7]", "FUN 3: TICK [8]", "FUN 3: TICK [9]", then "FUN 4 IN BURST[2197]", followed by "FUN 4: TICK [0]" through "FUN 4: TICK [9]". Then it shows "FUN 3 IN BURST[2229]", "FUN 3: TICK [0]", "FUN 3: TICK [1]", "FUN 3: TICK [2]", then "50ms time quantum ends", then "FUN 4 IN BURST[2198]", "FUN 4: TICK [0]", "FUN 4: TICK [1]", "FUN 4: TICK [1]", and finally "FUN 4: TICK [2]".

```
Bochs x86 emulator, http://
A: B: CD
FUN 3: TICK [7]
FUN 3: TICK [8]
FUN 3: TICK [9]
FUN 4 IN BURST[2197]
FUN 4: TICK [0]
FUN 4: TICK [1]
FUN 4: TICK [2]
FUN 4: TICK [3]
FUN 4: TICK [4]
FUN 4: TICK [5]
FUN 4: TICK [6]
FUN 4: TICK [7]
FUN 4: TICK [8]
FUN 4: TICK [8]
FUN 4: TICK [9]
FUN 3 IN BURST[2229]
FUN 3: TICK [0]
FUN 3: TICK [1]
FUN 3: TICK [2]
50ms time quantum ends
FUN 4 IN BURST[2198]
FUN 4: TICK [0]
FUN 4: TICK [1]
FUN 4: TICK [1]
FUN 4: TICK [2]
```



```
FUN 4: TICK [0]
50ms time quantum ends
FUN 3 IN BURST[2218]
FUN 3: TICK [0]
FUN 3: TICK [1]
FUN 3: TICK [2]
FUN 3: TICK [4]
FUN 3: TICK [5]
FUN 3: TICK [6]
FUN 3: TICK [7]
FUN 3: TICK [8]
FUN 3: TICK [9]

FUN 4: TICK [1]
FUN 4: TICK [2]
FUN 4: TICK [3]
FUN 4: TICK [4]
FUN 4: TICK [5]
FUN 4: TICK [6]
FUN 4: TICK [7]
FUN 4: TICK [7]
FUN 4: TICK [8]
FUN 4: TICK [9]
FUN 3 IN BURST[2219]
FUN 3: TICK [0]
```

